

WordleGym:

Benchmarking Wordle Strategies under Standard, Adversarial, and Hidden-Mode Play

Michael Wang

Math 242 • Duke University • Spring 2026

Code: github.com/Michael-OvO/WordleGym • Live demo: wordlegym.vercel.app

Problem Statement

Wordle is a word game in which the player has six guesses to identify a hidden five-letter word. After each guess, every letter is marked green (correct letter in the correct position), yellow (correct letter in the wrong position), or gray (not in the word). From this simple feedback rule emerges a rich decision problem: with 2315 possible answers, which guess at each turn makes the best use of the remaining information?

To compare strategies rigorously, I built *WordleGym*, a Python benchmark that plays the game automatically under different policies and records their performance. I studied three variants of the game in order to make the analysis more interesting:

Standard is the ordinary game: the answer is fixed at the start, and feedback is truthful.

Evil is an adversarial version. The game does not commit to a hidden answer in advance; instead, after each guess it returns the feedback pattern that preserves the largest number of still-valid answers. The adversary is cheating, but in a consistent and deterministic way.

Unknown combines the two. At the start of each game the game is secretly either Standard or Evil with equal probability, and the player must infer which from feedback alone.

The central question of the project is how simple, one-step strategies compare to one another, and how close they come to the best possible play.

Setting Up the Math

Some notation is needed before introducing the strategies. At any point in a game, let C denote the set of words still consistent with the feedback observed so far. A guess g partitions C into *buckets*: all words in C that would produce the same feedback pattern against g fall into the same bucket. Write B_r for the bucket corresponding to feedback pattern r , and $|B_r|$ for the number of words in it.

Intuitively, a good guess splits C into many small buckets. If one bucket is large, the guess barely narrows the search; if every bucket contains only one or two words, the guess is highly informative. This notion of “how evenly does a guess split the candidates” is the common thread linking all of the strategies I tested.

For each of the three modes, the optimal policy can be written as a recursion. These recurrences define what a perfect player would do; every strategy in the benchmark is an approximation to one of them.

In Standard mode, the minimum expected number of remaining guesses from a candidate set C satisfies

$$V(C) = \min_g \left[1 + \sum_{r \neq \gamma, B_r \neq \emptyset} \frac{|B_r|}{|C|} V(B_r) \right], \quad V(\{w\}) = 1. \quad (1)$$

Here γ denotes the all-green feedback pattern, which terminates the game immediately and so contributes no future cost. The interpretation is: choose the guess g that, averaged over where the true answer might lie, minimizes the number of guesses remaining.

In Evil mode, the adversary always returns the largest bucket, so the player wants to minimize the size of the bucket returned. Letting $T(C, g)$ denote that adversarial successor bucket,

$$D(C) = \min_g [1 + D(T(C, g))]. \quad (2)$$

In Unknown mode, the player maintains a posterior probability q that the current game is Standard, updated by Bayes' rule after each feedback. The optimal policy is a weighted combination of the Standard and Evil recurrences:

$$V_U(C) = \min_g [1 + q \cdot (\text{standard term}) + (1-q) \cdot (\text{evil term})]. \quad (3)$$

These recurrences are the exact optima, but solving them directly requires searching the entire game tree, which is computationally prohibitive in pure Python. The benchmark question then becomes: how well can more tractable, one-step strategies approximate these optima?

Strategies Tested

I implemented ten strategies, organized in tiers so that the results admit an apples-to-apples comparison.

Baselines. Two unsophisticated policies set a floor. **random-valid** chooses uniformly at random among the still-valid words, and **letter-frequency** chooses the word whose letters occur most frequently among the remaining candidates. Neither policy reasons about feedback patterns.

One-step heuristics reflect the conventional advice seen in popular Wordle write-ups, recast as formal optimization problems over a single guess. Each of these examines every allowed guess and selects the one that best satisfies a specific local objective:

- **expected-entropy** selects the guess with the largest Shannon *information gain*, which corresponds to the most balanced distribution of feedback buckets.
- **candidate-elim.** selects the guess that minimizes the expected number of candidates remaining after feedback is received.
- **minimax** is conservative: it selects the guess whose *largest* bucket is smallest, guarding against the worst-case feedback.

All three policies are intuitively close to optimal, but they are greedy — each reasons only one turn ahead.

Mode-aware strategies attempt to do better in Unknown mode by blending the Standard and Evil objectives according to the current posterior q . **posterior-hybrid** blends normalized entropy with evil-bucket reduction, weighted by q . **posterior-expectimax** uses a cleaner blend,

$$q \sum_r \frac{|B_r|^2}{|C|} + (1-q) |T(C, g)|,$$

which reduces exactly to **candidate-elim.** when $q = 1$ and to a pure evil-bucket minimizer when $q = 0$. **robust-scalarization** sidesteps the posterior entirely and picks the guess whose worst-case cost across both modes is smallest, minimizing $\max(\sum_r |B_r|^2/|C|, |T(C, g)|)$. In addition, **evil-shortest-path** is a single-ply version of the Evil DP, picking the guess that minimizes the adversary’s next bucket; it is the natural greedy approximation to **evil-dp** and works in all three modes.

Evil DP (exact). For Evil mode, I was able to solve recurrence (2) exactly. Because the evil adversary is deterministic, the reachable game tree contains only a few thousand distinct states, which is tractable in pure Python. The **evil-dp** strategy performs memoized depth-first search over this tree and returns the true optimum. Details of the algorithm are given below.

Benchmark Results

Every strategy played each of the 2315 canonical Wordle answers in Standard and Unknown mode. Evil mode records a single deterministic playthrough per strategy, since the adversary makes the outcome a fixed function of the policy. “Mean” is the average number of guesses, and “WC” is the worst case observed across the pool. The best value in each column is in **bold**.

Strategy	Standard		Evil		Unknown	
	Mean	WC	Mean	WC	Mean	WC
expected-entropy	3.465	6	5	5	4.232	6
posterior-hybrid	3.480	5	5	5	4.237	6
posterior-expectimax	3.485	5	5	5	4.248	5
candidate-elim.	3.486	5	5	5	4.243	5
evil-shortest-path	3.514	5	5	5	4.257	5
robust-scalarization	3.517	5	5	5	4.259	5
minimax	3.573	6	5	5	4.287	6
letter-frequency	3.587	8	5	5	4.294	8
random-valid	4.124	9	7	7	4.996	9
<i>evil-dp (optimal)</i>	—	—	4	4	—	—
Published optimum (Bertsimas & Paskov, 2022; Selby, 2022)	3.421	5	4	4	—	—

The same data is more striking as a scatter of average against worst-case guesses in Standard mode (Figure 1). The Pareto frontier of policies runs along the lower-left edge, where small gains in average performance come at the cost of larger worst cases.

Three observations from this table are worth emphasizing.

First, entropy is a surprisingly strong heuristic. The best one-step policy, **expected-entropy**, averages 3.465 guesses in Standard mode, while the published optimum reported by Bertsimas & Paskov (2022) is 3.421 — a gap of only 0.04 guesses, or roughly 1.3% above the optimum. A strategy that reasons only a single turn ahead is within a rounding error of the exact dynamic-programming optimum.

Second, average and worst-case performance are genuinely different objectives. **expected-entropy** has the best Standard-mode average but permits occasional six-guess games, whereas several strategies with slightly worse averages (such as **candidate-elim.** and **posterior-expectimax**) cap their worst case at 5. A strategy that minimizes one of these quantities does not, in general, minimize the other.

Third, the adversary is less powerful than intuition suggests. Every non-random strategy solves Evil mode in at most 5 guesses, and the exact DP solver (**evil-dp**) solves it in 4, matching the known




Figure 1: Average vs. worst-case guesses by strategy in Standard mode. The black star marks the published DP optimum of Bertsimas & Paskov (2022); Selby (2022). Points closer to the lower-left corner are strictly better.

optimum. Notably, $D(A) = 4$ is strictly less than the Standard-mode worst case of $W(A) = 5$: there exist specific fixed Wordle answers that are harder than anything the adversary can construct. The Evil tie-breaking rule, which forces the adversary to prefer the largest single bucket, removes enough flexibility that a well-chosen guess sequence can always corner the game.

The Evil DP

The Standard-mode optimum defined by (1) is too large to compute directly in pure Python. Each guess can produce up to 243 feedback patterns, and the reachable state space is on the order of 10^6 to 10^7 subsets. Existing exact solvers, such as those of Bertsimas & Paskov (2022) and Selby (2022), rely on compiled languages and custom data structures, and take hours of computation even then.

Evil mode is fundamentally more tractable. Because the adversary’s response is deterministic under a fixed tie-breaking rule, each guess has exactly one successor state rather than up to 243. The reachable state space from the full 2315-word answer set consequently contains only a few thousand distinct subsets, which is small enough for an unoptimized Python solver.

The `evil-dp` strategy implements recurrence (2) directly as memoized depth-first search. Starting from the full answer set, the solver considers each allowed guess, computes the adversary’s successor bucket, and recursively determines the minimum remaining guesses from that bucket. Results are cached on the canonical sorted tuple of C so that each subproblem is solved only once. A beam of $K = 100$ highest-ranked guesses per state is used to bound the branching factor; on this corpus, the optimal opening guess always falls within this beam, so the beam does not sacrifice optimality in practice.

The solver determines that the optimal opening guess is **RAISE** and guarantees a solution in exactly 4 guesses, matching the published optimum $D(A) = 4$. The initial run takes approximately 83 seconds, after which the policy is cached to disk and later runs are instantaneous.

A subtle implementation issue was tie-breaking. When two buckets are equal in size, the adversary must choose one consistently, or else cached subproblem solutions become unreliable. I adopted the tie-breaking convention used by the online evil-Wordle community — largest bucket first, then fewer greens, then fewer yellows, then lexicographically smaller pattern — so that my results are directly

comparable to theirs.

Unknown Mode and the Posterior

In Unknown mode the player must infer the game type from feedback alone. I represent this belief as a running probability q that the current game is Standard, initialized at $q_0 = \frac{1}{2}$ and updated by Bayes' rule after each feedback signal.

In practice, this posterior concentrates on the true game type very quickly. Averaged across all games, the mean posterior on the correct mode is approximately 0.93 after one guess, 0.99 after two guesses, and essentially 1 thereafter (Figure 2). Standard and Evil behave sufficiently differently after a single feedback pattern that a small number of turns is typically enough to identify the mode.




Figure 2: Unknown-mode posterior concentration. Every non-random policy exceeds 0.99 confidence on the true game mode within two guesses; random play requires three to four.

The implication for strategy design is that mode-aware policies (**posterior-hybrid**, **posterior-expectimax**) do not actually outperform the mode-agnostic **expected-entropy** baseline in Unknown mode. The best Unknown-mode average is **expected-entropy** at 4.232 guesses; the best mode-aware policy, **posterior-hybrid**, trails at 4.237. Once q is close to 0 or 1 — which happens within two turns — all policies effectively agree on the best next guess, so mode-aware blending only matters during the brief window before the posterior commits. The benchmark suggests that window is too short for the blending to pay off on this corpus.

Discussion

Several findings from the benchmark were unexpected to me.

One-step policies are nearly optimal. I initially expected a larger gap between greedy heuristics and the true optimum. In this problem, however, a strategy that picks the single most informative next guess averages only 1.3% more guesses than the exact DP. Lookahead is not useless, but the marginal benefit is modest.

Average and worst-case metrics disagree. Among the top strategies, those that minimize expected guesses do not minimize the worst case, and vice versa. Which metric to prioritize depends on the application: minimizing the average is appropriate when many games are played, while minimizing the worst case is appropriate when a single-game guarantee is required.

Adversarial worst case is easier than Standard worst case. In Evil mode the exact optimum is $D(A) = 4$ guesses. In Standard mode the worst-case optimum is $W(A) = 5$: under optimal play, the hardest fixed Wordle answer still requires 5 guesses. The implication is that the deterministic evil adversary is actually a weaker opponent than the unluckiest possible fixed answer. Evil mode is only harder than Standard mode in expectation (Standard averages 3.421).

Limitations and future work. The main limitation of the benchmark is that I could not implement an exact Standard-mode DP in pure Python within a reasonable time budget. A natural next step would be to vectorize the partition scoring using NumPy, which should give a 5–10 $\times$ speedup and might make the full DP feasible overnight. A more ambitious next step would be to port the inner loop to a compiled language.

Conclusion

I constructed a benchmark harness for Wordle that evaluates ten strategies across three game variants and compares their results to the published optima. The central finding is that simple one-step heuristics perform remarkably well: **expected-entropy** achieves an average of 3.465 guesses in Standard mode, within 0.04 guesses of the optimum of Bertsimas & Paskov (2022), and the minimax-style heuristics attain the optimal worst case of 5 guesses reported by Selby (2022).

A second, perhaps more surprising finding concerns Unknown mode. I expected the mode-aware blends (**posterior-hybrid**, **posterior-expectimax**) to beat the mode-agnostic **expected-entropy**. They do not: the Bayesian posterior over game type concentrates above 0.99 within two turns, so mode-aware blending only has a brief window in which to matter, and empirically that window is too small to move the average. The best Unknown-mode policy on this corpus is mode-agnostic.

Evil mode is the one variant where I was able to solve the recurrence exactly in pure Python, because the adversary’s deterministic tie-breaking rule collapses the reachable state space to a few thousand nodes. My solver recovers the published bound $D(A) = 4$ with opening guess **RAISE**. The one-guess gap between the DP and every greedy heuristic should not be read as a general statement about the power of dynamic programming; it is a consequence of determinism-induced tractability in this one mode.

Compared to popular Wordle advice, these results suggest that the commonly recommended heuristics (entropy maximization and minimax splitting) are close enough to optimal that the residual gap is not meaningfully exploitable by a human player. The cost of greediness is real but small.

Finally, the benchmark’s response to modeling changes is itself informative. The information-theoretic lower bound on expected guesses grows only as $\log_2 N$, so expected depth is relatively insensitive to the size of the answer list. Conversely, if the evil adversary were redefined to prefer the smallest consistent bucket, Evil mode would collapse toward Standard mode and $D(A)$ would drop sharply. Both perturbations are a single configuration change away in the engine, which was one of the original motivations for developing WordleGym as a reusable benchmark rather than a one-off analysis script.

References

- Bertsimas, D., & Paskov, A. (2022). An exact and interpretable solution to Wordle. *MIT Sloan Operations Research Center*. Available at <https://mitsloan.mit.edu/ideas-made-to-matter/how-algorithm-solves-wordle>.
- Selby, A. (2022). The best strategies for Wordle. *sonorouschocolate.com*. Available at https://sonorouschocolate.com/notes/index.php?title=The_best_strategies_for_Wordle.

Shannon, C. E. (1948). A mathematical theory of communication. *The Bell System Technical Journal*, 27(3), 379–423.

Wardle, J. (2022). Wordle [game]. *The New York Times*. Canonical answer and allowed-guess lists at <https://www.nytimes.com/games/wordle>.